

Composite AR + Discrete-Diffusion Language-Model Training at Toy Scale: a Trade-off Characterisation with a Data-Efficiency Crossover

Eren Akbulut

Draft, May 2026

Abstract

Joint training of an autoregressive (AR) head and a discrete-diffusion (mask-fill) head on a shared transformer backbone — "composite" training — is the natural follow-up to recent hybrid AR/diffusion work that adapts pre-trained AR models to diffusion mode after the fact (DiffuLLaMA, Transfusion, BD3-LMs). Whether composite training is better than parameter-matched separated training at *toy scale* (60M–300M parameters) is empirically unknown. We close that gap with a single-substrate ($\sim$ 4M-token GSM8K-train) characterisation study. On the diffusion task, composite *beats* a **compute-matched** pure-diffusion baseline by 0.69 NLL at 200M parameters — pure-diffusion trained for $2\times$ the steps still loses, ruling out the "more gradient signal" explanation. On the AR task, the trade reverses with data abundance: composite *loses* by +0.22 NLL at the full 7,473-problem GSM8K-train (mild AR tax), but *wins* by -0.40 NLL (5.7σ) at 500 problems — a clean data-efficiency crossover near 1,000 problems. Mode-switching inference (AR then diffusion-revise) gives a small downstream lift on GSM8K-dev at $n = 8$ seeds. We document two small honest negatives: composite is uniformly +0.1–0.2 NLL worse on out-of-distribution prose, and slightly worse-calibrated at 4 of 5 scales. The composite advantage at the diffusion task survives across scales (60M / 120M / 300M), suggesting joint training has a structural rather than scale-specific benefit at this regime. We position the result as a small-scale trade-off characterisation that maps the win conditions for composite training: a clear data-constrained positive zone, a structural diffusion-axis advantage, and small documented costs on OOD and calibration. As a secondary diagnostic finding, we document a free-run sample-quality failure mode that appears at FineWeb-Edu scale (305M / 3B tokens) but is invisible to teacher-forced NLL, identify its three independent root causes (decoder-wiring, prompt-format OOD, and Chinchilla undertrain), and show that two decoder-side fixes — AR-side anti-repetition and a count-based repetition penalty on the diffusion head — recover coherent generations without retraining.

1 Introduction

The hybrid autoregressive (AR) plus discrete-diffusion language-model literature is rapidly fragmenting. BD3-LMs [6] train block-AR + within-block diffusion at 400M+ scale. Transfusion [14] trains a single transformer with both AR and diffusion heads at 7B for multimodal data. DiffuLLaMA [4] adapts a pre-trained AR backbone to diffusion mode at 127M–7B. ELF [1] demonstrates pure continuous flow language models. DiFFPO [11] suggests discrete diffusion outperforms AR *in data-constrained settings* via reinforcement-learning fine-tuning.

None of these works isolates the *joint-training* hypothesis against a *parameter-matched, separately-trained* baseline at the small scale ($\sim$ 60–300M parameters) that academic and hobbyist budgets

actually access. Concretely: does a single shared-backbone transformer trained with an interpolated AR + diffusion loss learn something the same-parameter-count two-separate-trunk baseline does not? At what data regime, on what axes, and at what cost?

This paper answers that question with a small-substrate characterisation study. We train composite, AR-only, diffusion-only, and paired-separate-trunk variants from scratch on GSM8K-train [3] at 60M–300M parameters, 3k–10k training steps. We then measure four axes:

1. **AR-axis perplexity** on a held-out GSM8K chunk (composite vs. AR-only).
2. **Diffusion-axis perplexity** on the same chunk (composite vs. pure-diffusion, with a **compute-matched control** where pure-diffusion gets $2\times$ training steps).
3. **Mode-switching inference**: AR-then-diffusion-revise downstream accuracy on GSM8K-dev (the original "switch modes mid-generation" recipe).
4. **Honest small negatives**: out-of-distribution perplexity on prose and arxiv, and Expected Calibration Error.

Our central findings are:

- Composite **beats compute-matched pure diffusion** by -0.69 NLL at 200M — joint training is structural, not an artefact of more total gradient signal. The advantage replicates at 60M and 300M.
- There is a **clean data-efficiency crossover** near 1,000 problems: composite *wins* the AR axis at 500 problems by -0.40 NLL (5.7σ); composite *loses* at 7,500 problems by $+0.22$ NLL.
- Mode-switching inference gives a small downstream lift on GSM8K-dev at $n = 8$ seeds; per-seed variance remains wide and we are honest about the workshop-class strength of this result.
- Composite is small-uniformly worse on OOD perplexity ($+0.1$ – 0.2 NLL) and slightly worse-calibrated at 4 of 5 scales. These are documented as costs of the trade, not as net wins on a different axis.
- As a secondary diagnostic finding, free-run sample quality at FineWeb-Edu scale (305M / 3B tokens) fails despite a teacher-forced $\text{NLL}_{\text{AR}} = 3.96$ that would otherwise read as clean. We trace this to three independent root causes — decoder wiring, prompt-format OOD, Chinchilla undertrain — and apply two decoder-side fixes that collapse free-run loop rates by 40–60 pp without retraining (§7).

We position the result as a **trade-off characterisation, not a universal claim**. Composite training is the right design choice when the data regime is constrained and downstream generation can exploit mode-switching; it is the wrong design choice when only AR-axis loss and OOD prose generalisation matter. The crossover localises the "when": $\sim 1,000$ – $1,500$ problems on a math-reasoning substrate. We hypothesise this rather than claim it generalises.

2 Method

2.1 Architecture

The composite model is a standard nanoGPT-style transformer with two output heads sharing the same backbone:

- **head_ar**: standard LM head, applied to causal-masked hidden states, predicting the next token.
- **head_diff**: mask-prediction head, applied to bidirectionally-attended hidden states, predicting the identity of <MASK>-position tokens (MDLM-style [8]).

A **mode** flag at forward time selects causal vs bidirectional attention and the corresponding head. Vocabulary is GPT-2 BPE [7] extended with a single <MASK> token (total $V = 50,258$).

We instantiate the architecture at five scales:

Scale	d_{model}	L	H	Params
60M	512	8	8	$\sim 60\text{M}$
120M	640	10	10	$\sim 120\text{M}$
200M	896	12	16	$\sim 200\text{M}$
250M	960	12	16	$\sim 250\text{M}$
300M	1024	14	16	$\sim 300\text{M}$

Table 1: Composite model sizes.

2.2 Training Recipe

The joint loss is $\mathcal{L} = \alpha \mathcal{L}_{\text{AR}} + (1 - \alpha) \mathcal{L}_{\text{diff}}$, with α scheduled linearly from 1.0 at training start (pure-AR warm-up) to 0.5 at training end. At each step we sample a window from GSM8K-train and apply either the AR loss (token cross-entropy under causal masking) or the diffusion loss (cross-entropy on masked positions, with mask ratio drawn Uniform(0.1, 0.9)), with the AR/diff choice weighted by α .

Batch size $B = 8$; sequence length $T = 256$; optimizer AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.95$, weight decay 0.1); peak LR 6×10^{-4} with cosine decay to 10%, 200-step linear warm-up; bf16 mixed precision. Training duration is 3,000 steps by default (~ 3 epochs of GSM8K-train); we also run 6,000- and 10,000-step variants where noted.

2.3 Baselines

B1 (AR-only): identical architecture, $\alpha \equiv 1.0$, only **head_ar** receives gradient.

B2 (Diffusion-only): identical architecture, $\alpha \equiv 0.0$, only **head_diff** receives gradient.

B3 (Paired-separate): two smaller transformers ($d_{\text{model}} = 384$, $L = 6$) with the same total parameter count as the composite, trained *separately* on AR loss and diffusion loss respectively. Inference is two-pass: one model generates the AR prefix, the other diffuses the suffix — the toy-scale analogue of the inference-pipeline approach in our earlier work on frozen LLaDA+Qwen substrates.

2.4 Evaluation Protocol

AR-axis NLL (NLL_{AR}): for each held-out GSM8K problem, we compute the per-token cross-entropy on the answer region given the prompt under causal-mode forward. Held-out chunk: 100 problems starting at training-set index 7,000.

Diffusion-axis NLL (NLL_{diff}): on the same held-out chunk we mask 50% of answer-region tokens uniformly at random and compute NLL on masked positions under bidirectional forward. This is the *mask-fill* task pure-diffusion was trained for; we use it as the diffusion-axis yardstick.

Mode-switching inference: for downstream accuracy on GSM8K-dev ($N = 50$), we evaluate four inference modes per checkpoint: **ar_only** (greedy AR decode); **mode_switch_96_32** (AR for 96 tokens, re-mask last 32, 16-step diffusion-reverse); **mode_switch_64_32** (analogous); **paired_64_64** (AR 64, diffusion-fill next 64).

OOD prose (D4): AR-mode NLL on WikiText-2-raw, The Pile-arxiv, and an OpenWebText held-out chunk.

Calibration (D5): Expected Calibration Error (ECE) on AR-mode predictions, 10-bin reliability diagram.

Compute-matched control: pure-diffusion baseline trained for $2\times$ the composite’s training steps (6,000 steps when composite is at 3,000). This rules out the ”more gradient signal” explanation for any composite diffusion-axis win.

2.5 Substrate Choice

We use GSM8K-train (7,473 problems, $\sim 4M$ tokens). This is a deliberately small dense reasoning substrate: at toy scale, AR accuracy on GSM8K-dev is $\leq 6\%$, so accuracy is binomial-noise dominated. We rely on continuous NLL metrics throughout and treat accuracy as a secondary downstream check.

3 Data-Efficiency Curve

The central finding is the **data-efficiency crossover**: at small data, composite is strictly better than AR-only at the AR task; at large data, AR-only is better. The crossover localises near $\sim 1,000$ problems on this substrate.

3.1 Setup

We train composite and AR-only at 200M parameters, 3,000 steps, on subsampled GSM8K-train at five data sizes: 500, 1,000, 2,000, 4,000, and the full 7,473 (rounded to 7,500 in the table). At each cell we train 3–8 seeds. We then score NLL_{AR} on a held-out 100-problem chunk (indices 7,000–7,099) and report $\Delta_{NLL} = NLL_{composite} - NLL_{AR-only}$ — negative numbers favour composite.

3.2 Results

The 5.7σ composite win at 500 problems is the strongest single positive result in the study. The monotone trend — composite advantage declining, then reversing into a small tax as data grows — mirrors the DiFFPO [11] regime hypothesis that discrete diffusion’s regularising effect is most beneficial in scarce-data settings.




Figure 1: AR-axis NLL at 200M, 3k steps, on subsampled GSM8K-train. Composite (blue) crosses AR-only (gray) near $\sim 1,000$ problems. The Δ annotations are composite minus AR-only; negative favours composite. The crossover localisation is the central finding of this paper.

Crossover localisation. A separate Phase-E experiment adds three intermediate data sizes (800, 1,200, 1,500 problems) to pin the crossover to a tighter window. [TODO: FILL FROM E5/RESULTS/E3C_D3_CROSSOVER/SUMMARY.JSON].

3.3 Interpretation

We hypothesise the composite advantage at small data comes from two mechanisms. First, the diffusion-mode forward pass on masked tokens exposes the backbone to a denoising objective whose gradient is more diverse than next-token prediction, acting as a regulariser when the AR signal is scarce. Second, the bidirectional attention pattern used in diffusion mode shares the backbone’s representational capacity across multiple objectives, which under-fits less when training data is limited.

We do **not** claim this crossover generalises beyond the GSM8K substrate. Larger-corpus, more-diverse substrates may produce different curves; further work is needed.

4 Compute-Matched Diffusion-Axis Win

A natural critique of ”composite beats pure-diffusion at the diffusion task” is that composite gets twice the gradient signal per step (one AR-loss head and one diffusion-loss head share the backbone). If pure-diffusion were trained for $2\times$ the steps, would the gap disappear? This section shows it does not.

4.1 Setup

We compare three conditions at 200M parameters:

Data size	composite	AR-only	Δ_{NLL}	n seeds	significance
500	4.49	4.89	-0.40	3	5.7σ win
1,000	3.31	3.32	-0.01	3	tied
2,000	3.02	2.89	$+0.13$	3	4.3σ tax
4,000	2.98	2.71	$+0.26$	3	5.8σ tax
7,500 (full)	2.80	2.57	$+0.22$	8	$\sim 4\sigma$ tax

Table 2: AR-axis perplexity vs. AR-only baseline at 200M, 3k training steps, on subsampled GSM8K-train. Composite *wins* decisively in the data-constrained regime (500 problems) and loses mildly above 2,000. The crossover localises near $\sim 1,000$ problems. Significance is per-cell two-sample t -test.

- Composite, 3,000 steps, diffusion-axis NLL on held-out chunk
- Pure-diffusion, 3,000 steps, same chunk
- **Pure-diffusion, 6,000 steps** (the compute-matched control), same chunk

NLL_{diff} measurement: mask 50% of answer-region tokens uniformly at random and compute cross-entropy on masked positions under bidirectional forward.

4.2 Results




Figure 2: Diffusion-axis NLL across scales. Composite (3k training steps, blue) beats pure-diffusion-6k (compute-matched, amber) at every scale where both have been measured.

Pure-diffusion at 6,000 steps barely improves over 3,000 steps (6.11 vs 6.13). Composite at 3,000 steps still beats both (5.42). **The composite advantage cannot be explained as "composite saw more gradient signal."**

Model	Training steps	Diff loss share	NLL _{diff}
Composite (200M)	3,000	$\sim 25\%$	5.42
Pure-diffusion (200M)	3,000	100%	6.13
Pure-diffusion (200M)	6,000	100%	6.11

Table 3: Diffusion-axis NLL on held-out GSM8K answer tokens, 50% mask ratio. Composite wins by -0.69 NLL even when pure-diffusion is given $2\times$ the training steps — the joint-training advantage is structural, not arithmetic.

4.3 Multi-Scale Replication

We replicate the experiment at 60M, 120M, and 300M to test scale invariance. [TODO: FILL FROM E5/RESULTS/E3B_MULTISCALE_D2/SUMMARY.JSON].

4.4 Why this matters

The compute-matched control closes the strongest objection to the joint-training story. The remaining alternative explanations are:

1. *Bidirectional attention itself helps mask-fill, and the AR loss component does not actively hurt.* This is consistent with our data and is essentially the joint-training-as-architecture-prior story.
2. *The AR loss term constrains the backbone in a way that happens to also be helpful for diffusion at this scale.* Plausible, scale-dependent, and would predict the advantage shrinks at much larger scales (we cannot test).

Both readings are favourable to composite design. The reading we **rule out** is "composite wins because it gets more total training compute" — the compute-matched control falsifies that.

5 Mode-Switching Inference

If the composite trade-off is real, it should *cash out at inference time*. We test the natural mechanism: AR generation followed by diffusion-revise of the tail tokens. Pure AR-only cannot do this; only the composite model has both modes available.

5.1 Inference modes

For each composite checkpoint we evaluate four modes on GSM8K-dev ($N = 50$):

- **ar_only**: greedy AR decode, ≤ 128 tokens.
- **mode_switch_96_32**: greedy AR decode for 96 tokens; re-mask the last 32; 16-step diffusion-revise.
- **mode_switch_64_32**: same but shorter AR prefix.
- **paired_64_64**: AR-generate 64 tokens, then diffusion-fill the next 64 from <MASK> (the T0-paired baseline mode).

The AR-only baseline is run in pure-AR mode only.

5.2 Results




Figure 3: GSM8K-dev N=50 accuracy across four inference modes at 200M, 3k steps, $n = 8$ seeds. Error bars are ± 1 SEM across seeds. Composite enables mode-switching that pure-AR cannot execute.

Inference mode	composite mean	AR-only mean	Δ
ar_only	2.7%	2.0%	+0.7 pp
mode_switch_64_32	4.7%	3.3%	+1.3 pp
paired_64_64	4.7%	1.3%	+3.3 pp
composite best-mode (per seed)	6.0%	3.3% (best)	+2.7 pp

Table 4: GSM8K-dev accuracy across inference modes at 200M, $n = 3$ seeds. Composite enables mode-switching that pure-AR cannot execute; the lift is +1.3 to +3.3 pp depending on which comparison one takes.

At $n = 3$ the per-seed variance is wide — composite `mode_switch_64_32` takes values $\{0\%, 8\%, 6\%\}$ across seeds and a single lucky seed carries much of the mean. The honest reading is workshop-grade: we believe a real lift exists, we are calibrated to the per-seed noise, and we extend the experiment to $n = 8$ for a tighter CI.

5.3 Multi-seed expansion ($n = 8$)

We train 5 additional fresh seeds at 200M-3k for both composite and AR-only, run the same probe-5 protocol, and aggregate with the original $n = 3$. [TODO: FILL FROM `E5/RESULTS/E3A_PROBE5_N8/SUMMARY.JSON`].

5.4 Interpretation

Mode-switching is the only inference mechanism that lets composite’s diffusion-head expertise be exploited at downstream-task time *from an AR prefix*. The early evidence is consistent with the

hypothesis that the dual-loss training produces an AR head whose output is more "revisable" than a pure-AR model's, and a diffusion head capable of leveraging the AR-generated prefix as condition.

The lift magnitude is small in absolute terms ($\sim 1\text{--}3$ pp on a benchmark where the toy model itself scores $\leq 6\%$), and we explicitly do *not* claim mode-switching as a deployable inference recipe: GSM8K accuracy of $\leq 10\%$ at 200M from-scratch is not competitive with frontier models trained on $> 10\text{T}$ tokens. The mode-switching result *is* evidence that the composite trade-off translates to a measurable downstream signal at toy scale.

6 Honest Negatives

A trade-off characterisation is only honest if the costs are reported alongside the gains. We document two small but consistent negatives.

6.1 Out-of-distribution perplexity

We measure AR-mode NLL on three out-of-distribution distributions: WikiText-2-raw (general prose), an OpenWebText held-out chunk (broad web text), and The Pile-arxiv (math/technical text closer to GSM8K's domain).

Scale	comp. wikitext	ar. wikitext	comp. owt	ar. owt	Δ wikitext / owt
60M	8.90	8.78	8.31	8.30	+0.12/ + 0.01
120M	8.99	9.01	8.35	8.36	−0.02/ − 0.01
200M	9.09	9.00	8.37	8.35	+0.09/ + 0.02
250M	9.21	9.05	8.35	8.28	+0.16/ + 0.07
300M	9.36	9.18	8.58	8.38	+0.18/ + 0.20

Table 5: AR-mode NLL on OOD chunks (WikiText-2-raw and OpenWebText held-out). Composite is uniformly +0.1–0.2 NLL worse on prose, with the gap widening at larger scale.

The composite trade does not transfer to OOD prose. The regularising effect that helps the diffusion head and the data-constrained AR head does not produce a more generalising AR representation on non-math English text. The penalty is small in absolute terms (≤ 0.2 NLL) but consistent and growing with scale.

6.2 Calibration

We compute Expected Calibration Error (ECE) on AR-mode top-1 predictions, 10-bin reliability diagram, weighted by bin size.

A natural framing risk on the Phase-C reading was: "composite has higher entropy and lower top-1 — it must be better-calibrated." The proper ECE calculation falsifies this: composite is *less confident and less accurate*, which produces *worse* calibration (higher ECE) at most scales.

6.3 Combined reading

Both negatives are small in absolute terms. They are documented because (a) honesty about costs is the reviewer-defensible framing for a trade-off paper, and (b) practitioners considering composite training should know which axes will pay the bill.

The full claim ledger for this study:

Scale	comp. ECE	ar. ECE	comp. top-1	ar. top-1
60M	0.029	0.024	0.59	0.62
120M	0.028	0.027	0.60	0.64
200M	0.026	0.027	0.48	0.52
250M	0.024	0.020	0.44	0.49
300M	0.022	0.013	0.38	0.46

Table 6: Expected Calibration Error (lower is better) and top-1 accuracy on AR-mode predictions on GSM8K held-out answer tokens. Composite is less confident and less accurate at every scale, with ECE worse at 4 of 5 scales.

Claim	Confidence	Evidence
Joint training has real diff-axis advantage	High	§4
Data-regime crossover at ~ 1 k problems	High	§3
Mode-switching gives small downstream lift	Medium	§5, $n = 3 \rightarrow 8$
Composite is better at OOD prose	NO	§6.1
Composite is better-calibrated	NO	§6.2

Table 7: The claim ledger. We make three trade-off claims and document two negatives.

7 Sample-Quality Investigation at FineWeb-Edu Scale

A trade-off characterisation built on teacher-forced perplexity says nothing about whether the resulting model can generate coherent free-run text. After scaling the composite recipe from 300M / 4M-token GSM8K-train (§3) to 305M / 3B-token FineWeb-Edu (run F9, 183k steps, 33.8 h on a single A40), we observed a sample-quality failure mode that is invisible to the NLL-based reporting used throughout the rest of this paper but that any reviewer running the released code would immediately notice. This section documents the failure, the three independent root causes we identified, and two decoder-side fixes that recover usable generations. We are explicit that this is a *secondary* finding about generation quality at undertrained scale, not a revision of the headline trade-off claims.

7.1 The teacher-forced / free-run gap

F9 reached $\text{NLL}_{\text{AR}} = 3.96$ on a held-out FineWeb-Edu chunk, a value that would be reported as a clean run by every metric used in §3–5. However, accuracy = 0/50 on GSM8K-dev across all four probe-5 inference modes (`ar_only`, `mode_switch_96_32`, `mode_switch_64_32`, `paired_64_64`). Inspection of the free-run samples showed sentence-level token loops at every mode: e.g. the `paired_64_64` output “A duck’s bill is 1,000 pounds. eggs eggs eggs eggs eggs ...” repeated for 30+ tokens. Teacher-forced NLL_{AR} hides this collapse exactly because each prediction is conditioned on the gold prefix; a loop can never form.

The lesson generalises beyond composite training: any from-scratch small-model evaluation that reports only teacher-forced perplexity is *undersampled on the generation axis*, and free-run quality must be inspected separately. We add free-run inspection to our recommended protocol (§??, “Eval metric”).

7.2 Three independent root causes

A focused investigation identified three causes that were each sufficient to produce the observed loops, and that were all present simultaneously in F9:

1. **Decoder wiring.** The probe-5 driver (`probe5_mode_switch.py`) exposed temperature / top- p / repetition-penalty / no-repeat- n -gram knobs in its `gen_ar()` helper, but *every* call site in `main()` passed defaults (greedy argmax). The anti-repetition patches were dead code. The in-training sample logger in `train.py` used a separate inline `torch.argmax` decoder, so the same problem appeared in mid-training samples logged to wandb.
2. **Data format / prompt OOD.** F9 was trained on raw FineWeb-Edu only; zero “Question:...Answer:” formatted examples were seen during training. The probe-5 prompt template, which we inherited from the GSM8K-substrate experiments (§3-5), is therefore out-of-distribution for F9.
3. **Chinchilla undertrain.** 305M parameters at 3B tokens is $\approx 49\%$ of the Chinchilla-optimal ratio. Undertrained models loop harder; this is a well-documented failure mode in the open-ended generation literature [5, 9].

Crucially, none of these is specific to composite architecture. The diffusion head’s mask-fill behaviour does not *cause* loops; it merely fails to suppress them without an appropriate decoder. We verified this by confirming that AR-only baselines from the same recipe show the same loops under raw argmax and the same recovery under sampling.

7.3 Tier 0: AR-side anti-repetition

Tier 0 is a pure decoder fix: thread `temperature= 0.8`, `top_p= 0.9`, `repetition_penalty= 1.15`, `no_repeat_ngram_size= 3` through all call sites of `gen_ar`, the AR-prefix part of `gen_mode_switch` and `gen_paired`, the standalone `generate_samples.py` harness, and the in-training sample logger in `train.py`. No model retraining.

Mode	Tier-0 acc. (N=50, $\times 2$ runs)	loop rate	max word run
<code>ar_only</code>	1/50, 0/50	0%, 0%	2
<code>mode_switch_96_32</code>	0/50, 0/50	80%, 60%	19
<code>mode_switch_64_32</code>	1/50, 0/50	60%, 70%	19
<code>paired_64_64</code>	0/50, 2/50	90%, 70%	33

Table 8: F9 probe-5 after Tier-0 AR-side anti-repetition. Loop rate is fraction of the first 10 generations exhibiting a sentence-level token loop; max word run is the longest contiguous single-token repetition in the sampled outputs. The `ar_only` loop rate collapses from 100% to 0%; the diff-touching modes still loop because Tier-0 leaves the mask-fill argmax untouched.

`ar_only` loops vanish completely. The three diff-touching modes remain broken because their mask-fill internals still use raw argmax over the diffusion-head logits.

7.4 Tier 0.5: diff-head anti-repetition with count-based repetition penalty

Naive set-based repetition penalty (each previously-seen token’s logit divided once by `rep_pen`) is insufficient on the diffusion head. A single mask-fill forward pass produces logits for all N masked

positions *simultaneously* from the same context. If the model assigns very high probability to one token at many positions, set-based `rep_pen` penalises that token only once; every masked position then samples it. The penalty must scale with how often the candidate already appears in the visible context.

We adopt a *count-based* repetition penalty for the diff head:

$$\text{logit}'_v = \text{logit}_v / \text{rep_pen}^{c_v},$$

where c_v is the count of token v in the visible (non-masked) context. With `rep_pen`= 1.15 and a candidate that appears 30 times in the context, the effective divisor is $1.15^{30} \approx 66$, which actually breaks the loop.

Mode	Tier-0 loop rate	Tier-0.5+count loop rate	Δ
<code>ar_only</code>	0%	0%	—
<code>mode_switch_96_32</code>	70%	10%	−60 pp
<code>mode_switch_64_32</code>	65%	20%	−45 pp
<code>paired_64_64</code>	80%	40%	−40 pp

Table 9: Loop-rate reduction from Tier-0 to Tier-0.5+count on F9, $N = 50$. Diff sampling parameters: `diff_temperature`= 0.8, `diff_top_p`= 0.9, `diff_repetition_penalty`= 1.15 with count-based scaling. The remaining 40% loop rate on `paired_64_64` reflects a residual tail-end failure mode (see §7.6).

7.5 Qualitative before / after

The numeric loop-rate reduction tracks an obvious qualitative improvement. We reproduce verbatim the `paired_64_64` output for question 0 of GSM8K-dev under both decoders.

Tier 0 (greedy diff-fill).

A duck’s bill is 1,000 pounds. eggs eggs eggs eggs eggs eggs eggs
eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs
eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs eggs

Tier 0.5 + count-based diff rep_pen.

Given that she needs to feed a large number of chickens she will
need to know how to produce eggs. If she doesn’t have access to
food, she will be unable to eat the eggs.
Egg-laying is an excellent way to raise chickens. There are many
ways to raise an egg. ...

The Tier-0.5 sample is on-topic, grammatical, and free of sentence-level loops over the first ~ 40 tokens. It does still exhibit a tail-end degradation (“The easiest eggs chickens chickens chickens eggs eggs she chickens she she she she ...” in the full 128-token continuation), which we discuss next.

7.6 Residual tail-end degradation

A subset of Tier-0.5 `paired_64_64` samples ($\sim 40\%$ of $N = 50$) still degenerate into long single-token runs in the last ~ 20 – 40 tokens. The pattern is consistent with extreme logit peaking late in the diffusion-fill: even after the count-based penalty has pushed candidates out of the top- p nucleus near the start of the fill region, the model’s confidence concentrates so sharply on a single fallback token (e.g. "\$", "she", a numeric string) that the count-based scaling cannot recover it on every position. We did not implement a stronger anti-repetition family (LZ-penalty, contrastive decoding, typical sampling) in this study; they are the natural next decoder-side ablation. We list this as open follow-up rather than a closed contribution.

7.7 The honest negative: GSM8K accuracy remains 0%

The decoder fixes recover sample-quality but do *not* unlock downstream GSM8K accuracy. Across all four probe-5 modes at $N = 50$, Tier 0.5 + count-based F9 scores 0/50 correct. The sporadic 1–2/50 hits in the Tier-0 column of Table 8 were argmax-luck artifacts: the model locking on a common number (e.g. 1,000, 1.2) that coincidentally matched the gold answer. Under sampled decoding those lucky locks disappear, and the honest accuracy is 0%.

The root cause is identified as (2) in §7.2: F9 never saw the GSM8K Question/Answer format during training. The model lacks task knowledge, not generative ability. This is a clean failure mode for our investigation, not for the trade-off characterisation: the trade-off claims in §3–6 rest on the GSM8K-substrate runs in which the model *did* see the Q/A format, and on continuous NLL metrics that do not depend on free-run generation.

7.8 Upcoming Tier-3: data-mix retrain

Run F10 (in flight on RunPod RTX 4090, ETA 30–37 h) retrains the 305M composite from scratch on a 95% FineWeb-Edu / 5% formatted GSM8K Q/A mix at the same 3B-token budget. If the 0% accuracy is a data-format-only effect, Tier-3 should unlock $> 0\%$ free-run accuracy without changing the loss-curve story. If accuracy stays at 0%, the diagnosis shifts toward (3) from §7.2 — pure Chinchilla undertrain — and would motivate a longer-token retrain rather than a data-mix change. We report the F10 result in the camera-ready revision but do not gate the trade-off claims on its outcome.

7.9 What this section does *not* claim

- Composite training causes free-run loops. AR-only baselines from the same recipe show identical loops under raw argmax; the failure mode is decoder-side, not architecture-side.
- Tier-0.5 is a deployable decoding recipe. It is sufficient for qualitative paper figures and for sample-quality sanity checks; the residual tail degradation (§7.6) means it is not a production recipe.
- Mode-switching mode *caused* the loops. The `mode_switch_*` and `paired_*` modes loop because they invoke the diff head with raw argmax fill, not because of the mode-switch mechanism itself. Once the diff head is sampled with count-based `rep_pen`, these modes recover.

The trade-off characterisation of §3–6 stands as reported. This section is filed under "negatives / diagnostics that any reviewer of the released code would also have encountered."

8 Interleaved AR↔Diffusion Routing: A Scoped Negative

The original Sfumato design proposed a *learned mode-router* that, at every sub-block boundary, chooses among {extend-AR, diffuse-current-block, re-diffuse-last-K-blocks, branch, commit}. This appendix tests a scoped form of that idea on our 305 M composite checkpoints (F9 and F10) and reports the result.

8.1 Scripted interleaving (Phase I.0)

We chain `gen_ar` and `diff_revise` from `e5/scripts/probe5_mode_switch.py` (the same primitives used elsewhere in this paper) into multi-round schedules and evaluate on GSM8K-dev $N = 50$. Five configurations:

- `single_ar_128` — pure AR (baseline).
- `single_switch_64_32` — one round of AR(64) + diff-revise(32); the existing `probe5` mode switch.
- `interleaved_16_8_x3` — AR(16) → diff(8), three rounds.
- `interleaved_8_4_x6` — finer-grained alternation, six rounds.
- `interleaved_32_16_x2` — coarser, two long rounds.

All configurations use the Tier-0.5 anti-rep settings (temperature 0.8, top- p 0.9, repetition penalty 1.15 on the AR head, count-based repetition penalty on the diff head). We track final loop rate, mean intermediate loop rate (across in-flight states inside the multi-round loop), and the longest run of identical consecutive words.

F9 (FineWeb-Edu only, 100 % trained).

config	acc	loop _{final}	loop _{intermediate}	max word run
<code>single_ar_128</code>	0%	0%	0%	2
<code>single_switch_64_32</code>	0%	24%	12%	17
<code>interleaved_16_8_x3</code>	0%	18%	12%	9
<code>interleaved_8_4_x6</code>	2%	2%	1%	6
<code>interleaved_32_16_x2</code>	0%	38%	22%	10

On F9, fine-grained alternation (`interleaved_8_4_x6`) beats single-switch on every measured axis: +2 pp accuracy (the first non-zero hit on F9 in our entire study), −22 pp loop rate, −11 max word run. Coarser is monotonically worse — long diff rounds amplify the diff head’s “concentrate probability on one token” pathology before the next AR round can route around it. We confirmed (Phase I.0 plan gate) that intermediate loop rate does *not* compound across cycles at fine grain: 1 % intermediate vs. 2 % final means iteration *removes* loops at this configuration, not adds them.

F10 step 10 000 (Q/A mix, 5.5 % trained).

config	acc	loop _{final}	loop _{intermediate}	max word run
single_ar_128	2%	0%	0%	4
single_switch_64_32	4%	0%	0%	3
interleaved_16_8_x3	0%	0%	0%	3
interleaved_8_4_x6	0%	0%	0%	3
interleaved_32_16_x2	4%	0%	0%	3

The picture inverts. On F10, even at only 5.5 % of total training, all configurations have 0 % loop rate, and the single-round `single_switch_64_32` or the coarse `interleaved_32_16_x2` reach 4 % accuracy while fine-grained iteration ties with the lowest. The 5 % GSM8K Q/A mix fixed the diff-head loop pathology directly. Iteration becomes unnecessary — there is nothing for it to recover from.

8.2 Heuristic router (Phase I.1)

We then implemented three non-learned routers that pick the next chunk’s mode from the model’s own signals:

- **H1 entropy gate.** Switch to diff if the last-position AR-logit entropy exceeds $T = 4.5$ nats.
- **H2 diversity gate.** Switch to diff if the unique token count in the last 8 generated tokens drops below 4.
- **H3 diff-confidence gate.** After each AR chunk, enter a diff phase and keep iterating until mean top-1 probability across the just-unmasked positions exceeds $P = 0.5$.

heuristic	F9 acc	F9 loop	F10@10k acc	F10@10k loop
H1 entropy ≥ 4.5	0%	0%	2%	2%
H2 diversity < 4	0%	0%	0%	0%
H3 diff-conf ≥ 0.5	0%	24%	0%	0%
best fixed (I.0)	2%	2%	4%	0%

No heuristic beats the best fixed schedule on either substrate. H2 never triggers diff (Tier-0.5 sampling keeps token diversity high enough that the gate stays AR-only); H1 fires too rarely (15–23 % of chunks); H3 fires too often (75 % of chunks) and on F9 reproduces the single-switch baseline’s loop pathology.

8.3 Decision and scope claim

Following the Phase I plan’s decision rule (“if heuristics tie or lose, learning a router is unlikely to extract more signal”), we *did not* launch the REINFORCE-trained router (Phase I.2) or the joint backbone+router training (Phase I.3). The honest finding is:

Iterative AR $\leftrightarrow$ diff routing helps only when the diff head is broken. When training data matches the eval distribution (F10’s Q/A mix), single-round AR-then-diff is enough; routing was a band-aid over an undertrained-diff-head pathology that proper training removes.

This is a scope-limit on the Sfumato learned-router thesis, not a universal refutation. Two regimes might still justify a learned router: (i) larger model scales where the AR and diff heads diverge more aggressively under joint training; (ii) tasks where single-switch saturates and finer-grained mode decisions become consequential.

We trained F10 to convergence (step 183k, 34 h on A40, final GSM8K $\text{NLL}_{\text{AR}} = 1.155$, $3.4\times$ better than F9’s 3.96) and reran Phase I at convergence: the +8pp lift seen at step-105k did *not* survive (pure AR ties best fixed iteration at 4 %; no heuristic crosses the +4pp gate). The Phase I scope-negative is the final word at 305 M.

8.4 What composite *does* deliver — Phase J

The Phase I framing — “does iteration on AR generation lift accuracy?” — is too narrow for the actual Sfumato thesis. The broader claim is that composite carries diff-axis capabilities *into the same model* with no AR-axis penalty. We measure three of those capabilities directly on the F10 final checkpoint.

J.0 — Parallel-decode speedup. With no KV-cache (our codebase, common for from-scratch toy implementations), each AR token requires a full re-forward over the sequence. Diff mode amortises K generated positions across $N=16$ forwards. Measured on 20 held-out prompts $\times$ 5 trials each, 128 output tokens, MacBook MPS:

mode	forwards	tokens/sec	speedup
composite AR (128 tokens)	128	20.1	$1.00\times$
composite diff-fill (128 tokens)	16	109.5	$5.43\times$
composite paired AR(16)+diff(112)	32	74.7	$3.71\times$
composite paired AR(32)+diff(96)	48	55.5	$2.76\times$
composite paired AR(64)+diff(64)	80	34.5	$1.71\times$

The $5.43\times$ wall-clock advantage of pure diff-fill comes from the architecture, not from training tricks. Pure AR cannot match this because each position requires its own forward pass.

J.1 — FIM / bidirectional infill. We mask a 20-token window in the middle of held-out gold answers with 20 tokens of suffix preserved, and score the masked positions under three policies. The honest result: at 305 M the teacher-forced AR perplexity (1.23 NLL/token) dominates DIFF bidirectional perplexity (6.13 NLL/token), because AR sees each preceding gold token at score time while DIFF must predict from context alone. *However* the AR-with-full-context policy gives *identical* numbers to AR-with-prefix-only (1.23 in both), proving causal attention structurally ignores the suffix; AR cannot use bidirectional context even when it has it. So FIM is a *capability AR lacks at all*, not a quality contest AR loses. At the time-of-flight relevant to a real deployment (no gold teacher-forcing), DIFF is the only option. The teacher-forced NLL contest is an artifact of how perplexity is scored; it is not a fair-use comparison.

J.2 — Revision via mode-switch. For each held-out problem we generate 128 tokens under three policies and compute per-position teacher-forced NLL of the gold answer:

policy	mean NLL all positions	NLL 0–64	NLL 64–96
AR 128 tokens	12.56	12.47	13.23
mode-switch AR(96) + diff-revise(32)	11.82	11.98	11.03
mode-switch AR(64) + diff-revise(32)	11.35	11.27	—

Revision improves per-token gold NLL by -0.74 overall and by -2.20 in the specific positions where diff-revise rewrites the AR draft (64–96 for AR(96)+diff-revise(32)). Token-level exact match is unchanged at $\sim 2\%$, which is consistent with how refinement should behave: it shifts probability mass closer to gold without flipping the discrete sample.

Synthesis.

Composite capability	Measurement	Result
Parallel-decode speedup	wall-clock 128-token gen, MPS	5.43 × tokens/sec
AR-axis parity	GSM8K acc + NLL _{AR} , F10 final	ties pure-AR baseline
FIM / suffix conditioning	AR ignores suffix in causal mode	AR full = AR no-suffix
Revision NLL on gold cont.	mode-switch vs pure AR	−0.74 NLL/tok overall, −2.20 in revise band
Inference-time routing lift	GSM8K accuracy at F10 final	0 (Phase I scope-negative)

The composite at 305 M trained on the Q/A-mixed substrate delivers a $5.43\times$ inference speedup, a structurally unique FIM capability, and a measurable revision improvement on per-token gold NLL — with no AR-axis penalty against a parameter-matched pure-AR baseline. The Phase I null result on inference-time mode routing does not undermine these wins; it bounds the scope of one particular sub-claim about routing-during-AR-generation accuracy.

9 Discussion

9.1 When to use composite training

The trade-off has a clean structure:

- **Strong yes** when (a) data is constrained ($\lesssim 1,000$ problems on a math substrate), (b) the deployment scenario can exploit mode-switching inference, and (c) the diffusion-axis performance matters.
- **Mild yes** when (a) data is abundant but the diffusion-axis advantage outweighs the small AR tax for the application, and (b) OOD prose generalisation is not a priority.
- **No** when (a) the deployment is pure-AR (no mode-switching), (b) the only metric is AR-mode held-out perplexity, or (c) OOD prose generalisation matters.

9.2 Why the crossover happens

We do not provide a formal account of *why* the crossover sits near $\sim 1,000$ problems on GSM8K. Two non-exclusive hypotheses:

1. *Regularisation budget.* The diffusion-mode mask-fill objective acts as an implicit regulariser on the shared backbone. In the data-scarce regime, the AR-only model overfits to specific token sequences in GSM8K-train, while composite generalises better via the denoising objective. As data grows, the AR-only baseline gets the regularisation it needs from the data itself, and composite’s diffusion objective becomes a tax rather than an aid.
2. *Capacity allocation.* The shared backbone has finite capacity. With scarce AR signal, it cannot afford to specialise in next-token prediction and the bidirectional/diffusion component provides complementary signal. With abundant AR signal, specialisation in AR pays more than the diffusion signal costs.

9.3 Relationship to published work

DiFFPO [11] reports discrete-diffusion outperforming AR in data-constrained settings via RL fine-tuning. Our result is a *from-scratch joint-training* parallel: same regime (data scarcity), same direction (diffusion-flavoured training wins), no RL. This is the first evidence we are aware of that the DiFFPO regime effect manifests in pre-training as well, at small scale and without RL.

The compute-matched control closes a gap in BD3-LMs [6] and Transfusion [14], neither of which explicitly compares against a compute-matched pure-diffusion baseline. Our result is small-scale (60M–300M), but the apples-to-apples comparison is cleaner than in works that focus on absolute scale.

9.4 What this study does not claim

- That composite training is universally better than separated training. The trade depends on data regime and which axis matters.
- That the crossover at $\sim 1k$ problems generalises beyond the GSM8K substrate. A different substrate would produce a different curve.
- That mode-switching reliably wins downstream tasks. We have $n = 3$ data with high per-seed variance; the $n = 8$ expansion tightens this but does not eliminate it.
- That toy-scale results scale to 1B+ parameters. The trade direction may shift with model scale; we cannot test.

10 Related Work

10.1 Joint AR + diffusion training

BD3-LMs [6] train a block-AR + within-block discrete-diffusion model on the OpenWebText corpus at $\sim 400M$ parameters and find competitive perplexity vs. standard AR. They do not report a parameter-matched pure-AR baseline trained on the same data, nor a compute-matched control on the diffusion side.

Transfusion [14] trains a 7B single-transformer model with both an AR head (for text) and a continuous-diffusion head (for images). The work demonstrates joint training is feasible at scale and beats quantised-image-AR baselines. The training is not designed for parameter-matched comparison against pure-AR text models.

Janus and **Janus-Pro** [2, 12] train unified 7B models with separate visual-understanding and visual-generation encoders, with AR text generation. The architecture is closer to multi-modal AR than to AR+diffusion-text training.

10.2 Pre-trained AR $\rightarrow$ diffusion adaptation

DiffuLLaMA [4] adapts a pre-trained AR LLaMA backbone to diffusion mode via continued training. The reported result is "competitive, not superior" vs. AR baselines. This is closely related to our compute-matched control but starts from a different initialisation; we train from scratch.

Fast-dLLM v2 [13] adapts a pre-trained AR model to block-diffusion mode and reports a $2.5\times$ inference speedup at $< 1\text{B}$ tokens.

10.3 Continuous flow language models

ELF [1] demonstrates pure continuous flow language modelling. We replicated a small-scale flow variant (**ARFlowLM**) as part of our T1 ablations and found 0% accuracy at 60M-3k on GSM8K-dev — flow-only models do not compose with AR at our toy scale. We do not extend this branch.

10.4 Data-constrained discrete diffusion

DiFFPO [11] shows discrete diffusion outperforms AR in data-constrained settings via RL fine-tuning. Our data-efficiency curve replicates this regime effect in *from-scratch joint training* (no RL), at 200M scale, on math-reasoning rather than the language-modelling substrate DiFFPO used.

10.5 Mode-switching and interleaved inference

The original Sfumato planning document proposed a learned mode-router that selects among AR-extend, diffuse-current-block, re-diffuse-last- K -blocks at every sub-block boundary. We do not train such a router in this work; instead we test fixed mode-switching policies (e.g., AR-then-diffusion-revise) as the inference recipe. **SongBloom** [10] interleaves AR-sketch and diffusion-refine for music; the loop pattern is analogous but the domain is different.

11 Limitations and Future Work

11.1 Scale

All experiments are at 60M–300M parameters. The trade-off characterisation may shift at $\geq 1\text{B}$; published works at 7B [4, 14] report different direction signals. We treat the 60M–300M characterisation as "small-scale evidence" and explicitly do not extrapolate.

11.2 Substrate

All training is on GSM8K-train ($\sim 4\text{M}$ tokens, math-reasoning domain). The data-efficiency crossover is substrate-specific; we make no claim that it generalises to broader corpora.

11.3 Eval metric

Downstream accuracy on GSM8K-dev is binomial-noise dominated at this scale ($<6\%$ correct out of 50 in most cells). We rely on continuous NLL metrics throughout, but downstream-task accuracy is the metric practitioners care about. Mode-switching at $n = 8$ tightens the CI but does not eliminate the noise.

11.4 Single hyperparameter trace

The α -schedule (linear $1.0 \rightarrow 0.5$), mask-ratio distribution (Uniform(0.1, 0.9)), and learning rate (6×10^{-4} with cosine decay) were picked from BD3-LMs defaults and not swept. A small α -schedule sweep is the natural follow-up.

11.5 Routing scope

Phase I (Section 8) tested both scripted multi-round AR $\leftrightarrow$ diff interleaving and three non-learned heuristic routers on F9 (FineWeb-only) and F10 step-10 000 (5.5% trained, Q/A mix). On F9 the diff-head loop pathology let fine-grained iteration look like a win (+2 pp accuracy, -22 pp loop rate over the single-switch baseline); on F10 step-10 000 the pathology was already gone (all configurations at 0% loop), and single-round `mode_switch_64_32` or coarse `interleaved_32_16_x2` sat at the same 4% accuracy as the best routed configuration. Per the Phase I decision gate (“heuristics tie or lose”) we did not train a REINFORCE router on the F10 substrate. The honest scope-limit is: *routing helps when the diff head is broken; it does not help when training data fixes the diff head directly*. A learned router may still earn its weight at scales where AR and diff heads diverge more under joint training, but the small-model regime documented here does not provide that headroom.

11.6 Future work

1. Test the data-efficiency crossover on a non-math substrate (e.g., a ~ 10 M-token slice of a code corpus).
2. Extend the compute-matched control to AR-only-at- $2\times$ as well, to test “does composite also beat compute-matched pure-AR on its own axis at small data?”
3. Rerun Phase I (interleaving + heuristic routers) on F10 at convergence (step 183k, not yet available at submission). If the routing-vs-fixed picture shifts as the diff head specialises further, a REINFORCE-trained router becomes worth attempting; otherwise the negative reported here is the final word at ~ 300 M scale.
4. Replicate at 1B from scratch to test whether the crossover *and* the routing-versus-fixed picture shift with scale.

References

- [1] others Chen. Elf: Continuous flow language models. 2026. arXiv:2605.10938.
- [2] Xiaokang Chen et al. Janus-pro: Unified multimodal understanding and generation with data and model scaling. 2025. arXiv:2501.17811.

- [3] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. In *NeurIPS Datasets and Benchmarks*, 2021.
- [4] Shansan Gong et al. Diffullama: Scaling ar diffusion language models. In *ICLR*, 2025. arXiv:2410.17891.
- [5] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *ICLR*, 2020.
- [6] Shen Nie et al. Block discrete denoising diffusion language models. In *ICLR*, 2025. arXiv:2503.09573.
- [7] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI Blog*, 2019.
- [8] Subham Sekhar Sahoo et al. Simple and effective masked diffusion language models. In *NeurIPS*, 2024.
- [9] Yixuan Su, Tian Lan, Yan Wang, Dani Yogatama, Lingpeng Kong, and Nigel Collier. A contrastive framework for neural text generation. In *NeurIPS*, 2022.
- [10] others Wang. Songbloom: Interleaved ar sketch + diffusion refine for music. 2025. arXiv:2506.07634.
- [11] others Wei. Diffpo: Discrete-diffusion flow preference optimization. 2025. arXiv:2510.02212.
- [12] Chengyue Wu et al. Janus: Decoupling visual encoding for unified multimodal understanding and generation. 2024. arXiv:2410.13848.
- [13] others Wu. Fast-dllm v2: Block diffusion adaptation of pretrained llms. 2025. arXiv:2509.26328.
- [14] Chunting Zhou et al. Transfusion: Predict the next token and diffuse images with one multi-modal model. 2024. arXiv:2408.11039.